Coral

Get started with the USB Accelerator

The Coral USB Accelerator is a USB device that provides an Edge TPU as a coprocessor for your computer. It accelerates inferencing for your machine learning models when attached to either a Linux, Mac, or Windows host computer. This page is your guide to get started.

All you need to do is download the Edge TPU runtime and the PyCoral library on your computer. Then we'll show you how to run a TensorFlow Lite model using the accelerator.

If you want to learn more about the hardware, see the USB Accelerator datasheet.

Requirements

- A computer with one of the following operating systems:

  - Linux Debian 10, or a derivative thereof (such as Ubuntu 18.04), and a system architecture of either x86-64, Armv7 (32-bit), or Armv8 (64-bit) (Raspberry Pi is supported, but we have only tested Raspberry Pi 3 Model B+ and Raspberry Pi 4)

  - macOS 10.15, with either MacPorts or Homebrew installed

  - Windows 10

- One available USB port (for the best performance, use a USB 3.0 port)

- Python 3.5-3.8

1: Install the Edge TPU runtime

The Edge TPU runtime provides the core programming interface for the Edge TPU. You can install it on your host computer as follows, on Linux, on Mac, or on Windows.

1a: On Linux

1. Add our Debian package repository to your system:

```
echo "deb https://packages.cloud.google.com/apt coral-edgetpu-stable main" | sudo tee /etc

curl https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo apt-key add -

sudo apt-get update
```

2. Install the Edge TPU runtime:

```
sudo apt-get install libedgetpu1-std
```

3. Now connect the USB Accelerator to your computer using the provided USB 3.0 cable. If you already plugged it in, remove it and replug it so the newly-installed udev rule can take effect.

Then continue to install the PyCoral library.

Install with maximum operating frequency (optional)

The above command installs the standard Edge TPU runtime for Linux, which operates the device at a reduced clock frequency. You can instead install a runtime version that operates at the maximum clock frequency. This increases the inferencing speed but also increases power consumption and causes the USB Accelerator to become very hot.

If you're not certain your application requires increased performance, you should use the reduced operating frequency. Otherwise, you can install the maximum frequency runtime as follows:

```
sudo apt-get install libedgetpu1-max
```

You cannot have both versions of the runtime installed at the same time, but you can switch by simply installing the alternate runtime as shown above.

> **Caution:** When operating the device using the maximum clock frequency, the metal on the USB Accelerator can become very hot to the touch. This might cause burn injuries. To avoid injury, either keep the device out of reach when operating it at maximum frequency, or use the reduced clock frequency.

1b: On Mac

1. Download and unpack the Edge TPU runtime:

```
curl -O https://dl.google.com/coral/edgetpu_api/edgetpu_runtime_20210119.zip

unzip edgetpu_runtime_20210119.zip
```

2. Install the Edge TPU runtime:

```
cd edgetpu_runtime

sudo bash install.sh
```

The installation script will ask whether you want to enable the maximum operating frequency. Running at the maximum operating frequency increases the inferencing speed but also increases power consumption and causes the USB Accelerator to become very hot. If you're not certain your application requires increased performance, you should type "N" to use the reduced operating frequency. You can change this later by re-running this script.

You can read more about the performance setting in section 4.1 of the USB Accelerator datasheet.

3. Now connect the USB Accelerator to your computer using the provided USB 3.0 cable.

Then continue to install the PyCoral library.

1c: On Windows

1. First, make sure you have the latest version of the Microsoft Visual C++ 2019 redistributable.

2. Then download edgetpu_runtime_20210119.zip.

3. Extract the ZIP files and double-click the `install.bat` file inside.

A console opens to run the install script and it asks whether you want to enable the maximum operating frequency. Running at the maximum operating frequency increases the inferencing speed but also increases power consumption and causes the USB Accelerator to become very hot. If you're not certain your application requires increased performance, you should type "N" to use the reduced operating frequency. You can change this later by re-running this script.

You can read more about the performance setting in section 4.1 of the USB Accelerator datasheet.

4. Now connect the USB Accelerator to your computer using the provided USB 3.0 cable.

2: Install the PyCoral library

PyCoral is a Python library built on top of the TensorFlow Lite library to speed up your development and provide extra functionality for the Edge TPU.

We recommend you start with the PyCoral API, and we use this API in our example code below, because it simplifies the amount of code you must write to run an inference. But you can build your own projects using TensorFlow Lite directly, in either Python or C++.

To install the PyCoral library, use the following commands based on your system.

2a: On Linux

If you're using Debian-based Linux system (including a Raspberry Pi), install PyCoral (and TensorFlow Lite) as follows:

```
sudo apt-get install python3-pycoral
```

2b: On Mac and Windows

If you're using Mac or Windows, install PyCoral (and TensorFlow Lite) as follows:

```
pip3 install --extra-index-url https://google-coral.github.io/py-repo/ pycoral
```

3: Run a model on the Edge TPU

Now you're ready to run an inference on the Edge TPU.

Windows users: The following code relies on a Bash script to install dependencies. If you're new to using Bash on Windows, we suggest you try either Windows Subsystem for Linux or Git Bash from Git for Windows.

Follow these steps to perform image classification with our example code and MobileNet v2:

1. Download the example code from GitHub:

```
mkdir coral && cd coral

git clone https://github.com/google-coral/pycoral.git

cd pycoral
```

2. Download the model, labels, and bird photo:

```
bash examples/install_requirements.sh classify_image.py
```

3. Run the image classifier with the bird photo (shown in figure 1):

```
python3 examples/classify_image.py \
--model test_data/mobilenet_v2_1.0_224_inat_bird_quant_edgetpu.tflite \
--labels test_data/inat_bird_labels.txt \
--input test_data/parrot.jpg
```


Figure 1. parrot.jpg

You should see results like this:

```
----INFERENCE TIME----
Note: The first inference on Edge TPU is slow because it includes loading the model into E
11.8ms
3.0ms
2.8ms
2.9ms
2.9ms
-------RESULTS--------
Ara macao (Scarlet Macaw): 0.76562
```

Congrats! You just performed an inference on the Edge TPU using TensorFlow Lite.

To demonstrate varying inference speeds, the example repeats the same inference five times. Your inference speeds might differ based on your host system and whether you're using USB 2.0 or 3.0.

The top classification label is printed with the confidence score, from 0 to 1.0.

To learn more about how the code works, take a look at the `classify_image.py` **source code** and read about how to **run inference with TensorFlow Lite**.

> **Note:** The example above uses the PyCoral API, which calls into the TensorFlow Lite Python API, but you can instead directly call the TensorFlow Lite Python API or use the TensorFlow Lite C++ API. For more information about these options, read the **Edge TPU inferencing overview**.

Next steps

To run some other models, such as real-time object detection, pose estimation, keyphrase detection, on-device transfer learning, and others, check out our **example projects**. In particular, if you want to try running a model with camera input (including support for the Raspberry Pi camera), try one of the several **camera examples**.

If you want to train your own model, try these tutorials:

- **Retrain an image classification model using post-training quantization** (runs in Google Colab)

- **Retrain an image classification model using quantization-aware training** (runs in Docker)

- **Retrain an object detection model using quantization-aware training** (runs in Docker)

Or to create your own model that's compatible with the Edge TPU, read **TensorFlow Models on the Edge TPU**.